

CEG7450: Lecture 2

Reading

- **[Ste97]** W. Stevens, “TCP congestion control”, RFC 2001, Jan 1997.

Overview

➤ Layering

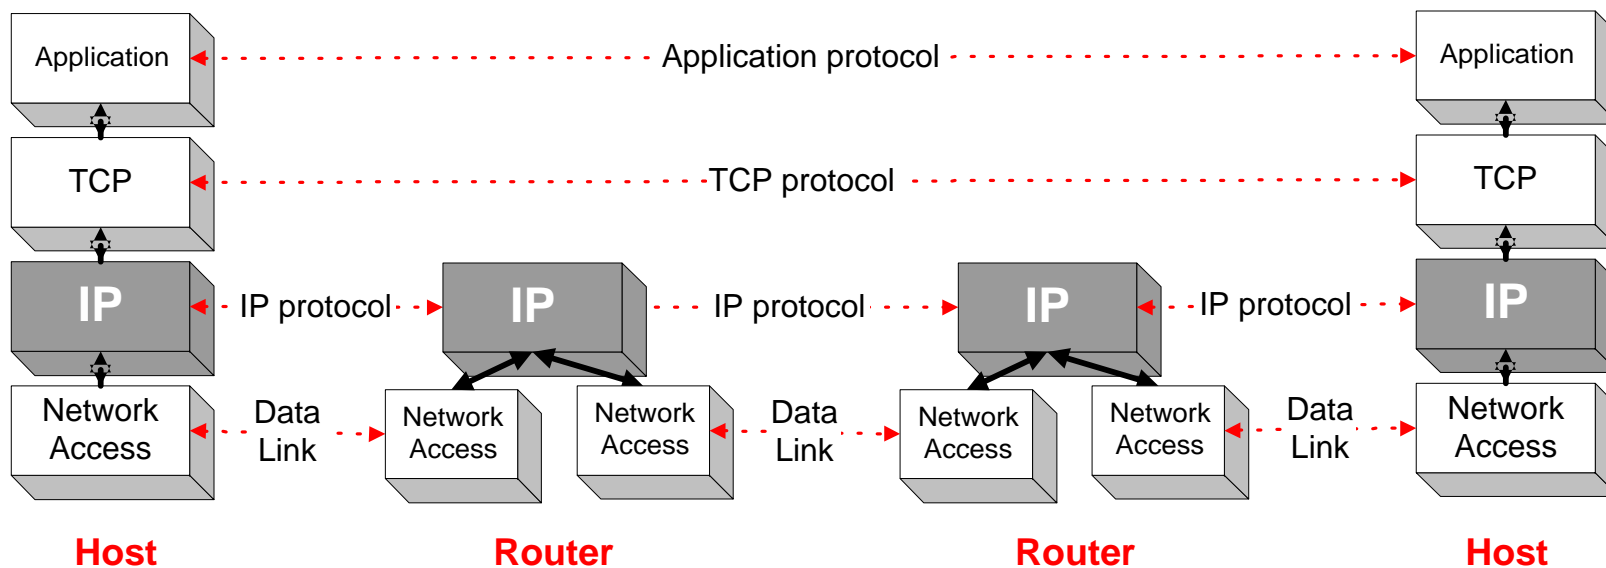
- End-to-End Arguments
- A Case Study: the Internet

What is Layering?


- A technique to organize a network system into a **succession** of logically distinct entities, such that the service provided by one entity is **solely** based on the service provided by the previous (lower level) entity

Network Protocols

- Specify any function that requires cooperation between two or more network entities
 - specify the format of the information that is sent/received among routers and end-systems
 - specify timings and the actions that a node has to take when it receives special messages or special events occur



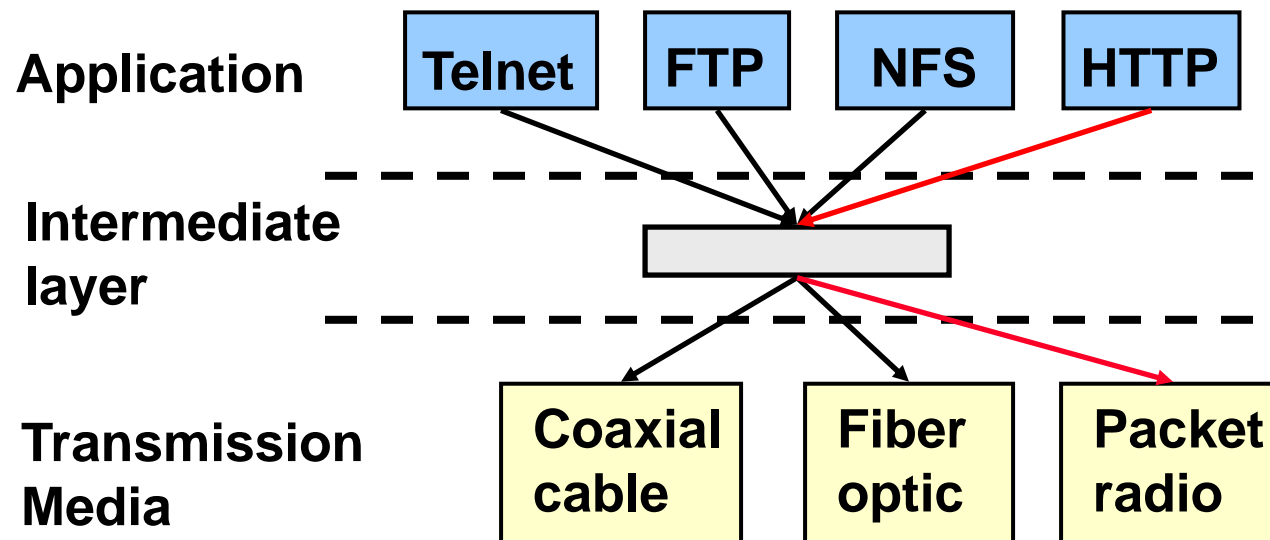
Why Layering?



- No layering: each new application has to be **re-**implemented for every network technology!

Why Layering?

- Solution: introduce an intermediate layer that provides a **unique** abstraction for various network technologies



Layering

- Advantages

- Modularity – protocols easier to manage and maintain
- Abstract functionality –lower layers can be changed **without** affecting the upper layers
- Reuse – upper layers can reuse the functionality provided by lower layers

- Disadvantages

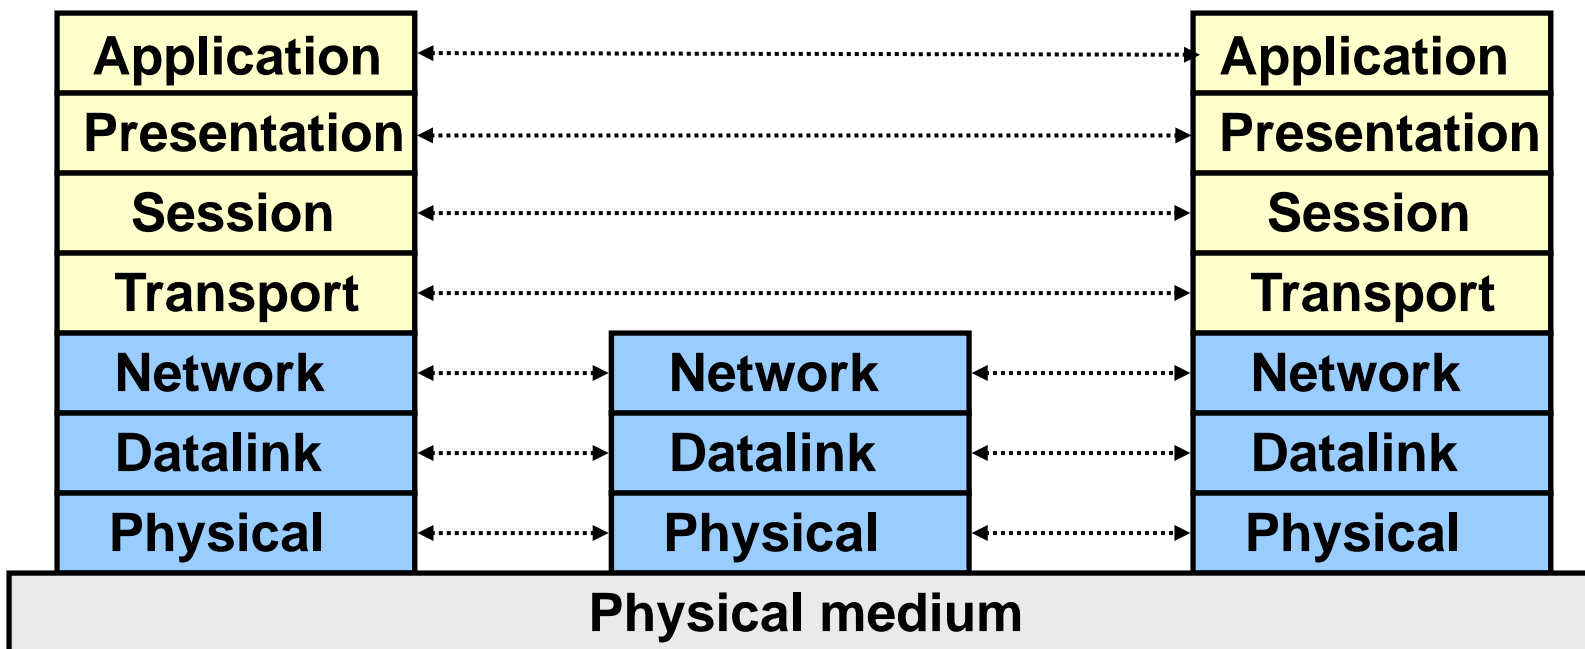
- Information hiding – inefficient implementations

ISO OSI Reference Model

- ISO – International Standard Organization
- OSI – Open System Interconnection
- Started in 1978; first standard 1979
 - ARPANET started in 1969; TCP/IP protocols ready by 1974
- Goal: a general **open** standard
 - allow vendors to enter the market by using their own implementation


ISO OSI Reference Model

- Seven layers
 - Lower three layers are peer-to-peer
 - Next four layers are end-to-end



Data Transmission

- A layer can use **only** the service provided by the layer immediate below it
- Each layer may change and add a header to data packet



OSI Model Concepts


- Service – says **what** a layer does
- Interface – says **how** to **access** the service
- Protocol – says **how** is the service **implemented**
 - a set of rules and formats that govern the communication between two peers

Physical Layer (1)

- Service: move the information between two systems connected by a physical link
- Interface: specifies how to send a bit
- Protocol: coding scheme used to represent a bit, voltage levels, duration of a bit

- Examples: coaxial cable, optical fiber links; transmitters, receivers

Encoding Schemes



- 4B/5B encode each 4 bits in a 5-bit code such that there is at most one leading 0 and at most two trailing zero and use NRZI encoding

Datalink Layer (2)

- Service:
 - framing, i.e., attach frames separator
 - send data frames between peers
 - others:
 - arbitrate the access to common physical media
 - ensure reliable transmission
 - provide flow control
- Interface: send a data unit (packet) to a machine connected to **same** physical media
- Protocol: layer addresses, implement Medium Access Control (MAC) (e.g., CSMA/CD)...

Marking the Beginning/Ending of a Frame

- Use a special bit sequence
- Problem: what happens if this sequence appears in the data payload?
- Use **bit stuffing** technique; e.g., assume both encoding, decoding bit sequences are 01111110
 - sender: inserts a 0 after five consecutive 1s
 - receiver: when it sees five 1s makes decision on next **two** bits
 - if next bit 0 (this is a stuffed bit), remove it
 - if next bit 1, look at the next bit
 - if 0 this is end-of-frame (receiver has seen 01111110)
 - if 1 this is an error, discard the frame (receiver has seen 01111111)

Network Layer (3)

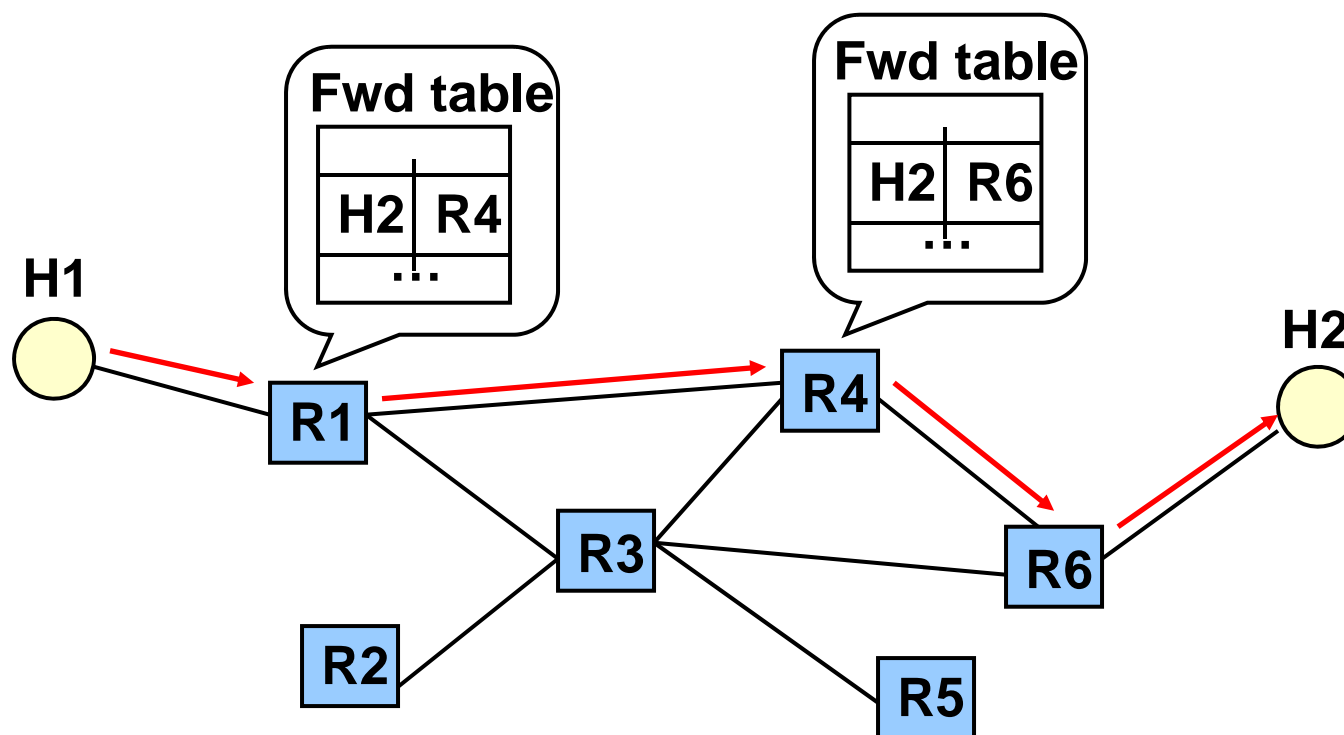
- Service:
 - deliver a packet to specified destination
 - perform segmentation/reassembly
 - others:
 - packet scheduling
 - buffer management
- Interface: send a packet to a specified destination
- Protocol: define global unique addresses;
construct routing tables

Data and Control Planes

- Data plane: concerned with
 - packet forwarding
 - buffer management
 - packet scheduling
- Control Plane: concerned with installing and maintaining state for data plane

Example: Routing

- Data plane: use Forwarding Table to forward packets
- Control plane: construct and maintain Forwarding Tables (e.g., Distance Vector, Link State protocols)



Transport Layer (4)

- Service:
 - provide an **error-free** and **flow-controlled** end-to-end connection
 - multiplex multiple transport connections to one network connection
 - split one transport connection in multiple network connections
- Interface: send a packet to specify destination
- Protocol: implement reliability and flow control
- Examples: TCP and UDP

Session Layer (5)

- Service:
 - full-duplex
 - access management, e.g., token control
 - synchronization, e.g., provide check points for long transfers
- Interface: depends on service
- Protocols: token management; insert checkpoints, implement roll-back functions

Presentation Layer (6)

- Service: convert data between various representations
- Interface: depends on service
- Protocol: define data formats, and rules to convert from one format to another

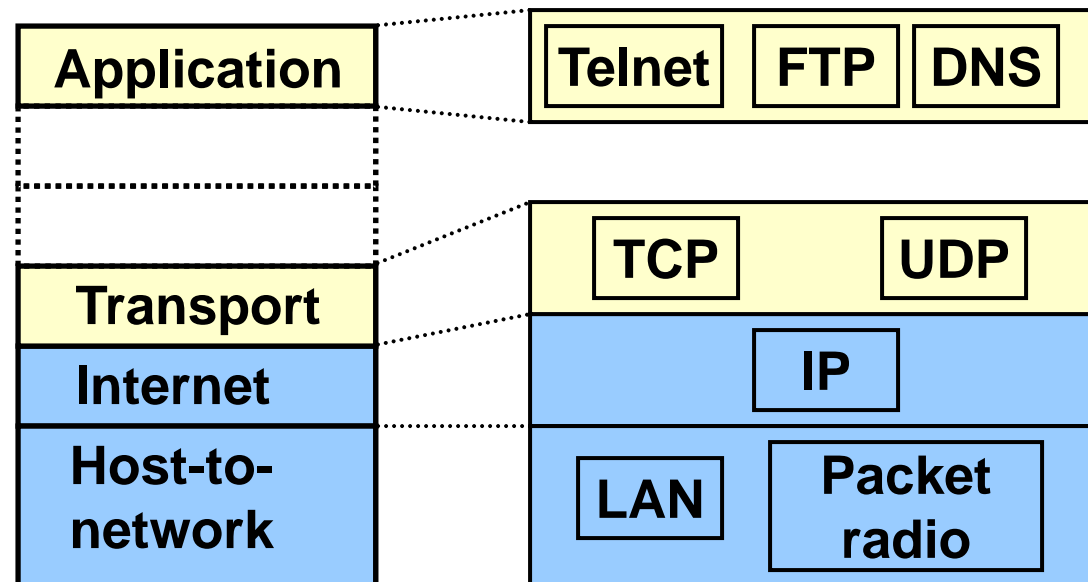

Application Layer (7)

- Service: any service provided to the end user
- Interface: depends on the application
- Protocol: depends on the application

- Examples: FTP, Telnet, WWW browser

OSI vs. TCP/IP

- OSI: conceptually define: service, interface, protocol
- Internet: provide a successful implementation



Key Design Decision

- How do you divide functionality across the layers?


Overview

- Layering
- End-to-End Arguments
- A Case Study: the Internet

End-to-End Argument

- Think twice before implementing a functionality that you believe that is useful to an application at a lower layer
- If the application can implement a functionality correctly, implement it at a lower layer **only** as a performance enhancement

Example: Reliable File Transfer



- Solution 1: make each step reliable, and then concatenate them
- Solution 2: end-to-end check and retry

Discussion

- Solution 1 not complete
 - What happens if the sender or/and receiver misbehave?
- The receiver has to do the check anyway!
- Thus, full functionality can be entirely implemented at application layer; **no** need for reliability from lower layers
- Is there any need to implement reliability at lower layers?

Discussion

- Yes, but only to improve performance
- Example:
 - assume a high error rate on communication network
 - then, a reliable communication service at data link layer might help

Trade-offs

- Application has more information about the data and the semantic of the service it requires (e.g., can check only at the end of each data unit)
- A lower layer has more information about constraints in data transmission (e.g., packet size, error rate)
- Note: these trade-offs are a direct result of layering!

Rule of Thumb

- Implementing a functionality at a lower level should have minimum performance impact on the application that do not use the functionality

Other Examples

- Secure transmission of data
- Duplicate message suppression
- Transaction management
- RISC vs. CISC

Overview


- Layering
- End-to-End Arguments
- A Case Study: the Internet

Internet & End-to-End Argument

- Network layer provides one simple service: best effort datagram (packet) delivery
- Only one higher level service implemented at transport layer: reliable data delivery (TCP)
 - performance enhancement; used by a large variety of applications (Telnet, FTP, HTTP)
 - does not impact other applications (can use UDP)
- Everything else implemented at application level

Internet Architecture

- Packet-switched datagram network
- IP is the glue
- Hourglass architecture
 - all hosts and routers run IP
- Stateless architecture
 - no per flow state inside network



Key Advantages

- The service can be implemented by a large variety of network technologies
- Does not require routers to maintain any fine grained state about traffic. Thus, network architecture is
 - Robust
 - Scalable
 - Stateless

What About Other Services?

- Multicast?
 - Multicast group management at routers?
 - End host based multicast?
- Quality of Service (QoS)?
 - Per flow state: resource reservation, classification, buffer management, scheduling
 - Stateless quality of service?

Summary: Layering

- Key technique to implement communication protocols; provides
 - Modularity
 - Abstraction
 - Reuse
- Key design decision: what functionality to put in each layer?


Summary: End-to-End Arguments

- If the application can do it, don't do it at a lower layer
 - the application knows the best what it needs anyway
 - add functionality in lower layers iff it is
 - (1) used and improves performance of a large number of applications, and
 - (2) does not hurt other applications
- Success story: Internet

Summary

- Challenge of building a good (network) system:
find the right balance between:

Reuse, implementation effort
(apply layering concepts)



- No universal answer: the answer depends on the goals and assumptions!