

CEG 7450: (TCP Congestion Control)

Reading

- **[BCS94]** R. Braden, D. Clark & S. Shenker. "Integrated Services in the Internet Architecture: an Overview", RFC 1633, June 1994.

2

Problem


- How much traffic do you send?
- Two components
 - flow control
 - make sure that the receiver can receive as fast as you send
 - congestion control
 - make sure that the network delivers the packets to the receiver
- However, in TCP, these mechanisms are inherently integrated with reliability

3

Go-Back-N

Sender:


- k-bit seq # in packet header
- "window" of up to N, consecutive unack'ed packets allowed



- timer for each in-flight packet
- *timeout(n)*: retransmit packet **n** and all higher seq # packets that were sent in window

4

Cumulative ACK




- ACK(n): ACKs all packets up to, including seq # n
 - may receive duplicate ACKs


5

Packet receiving order

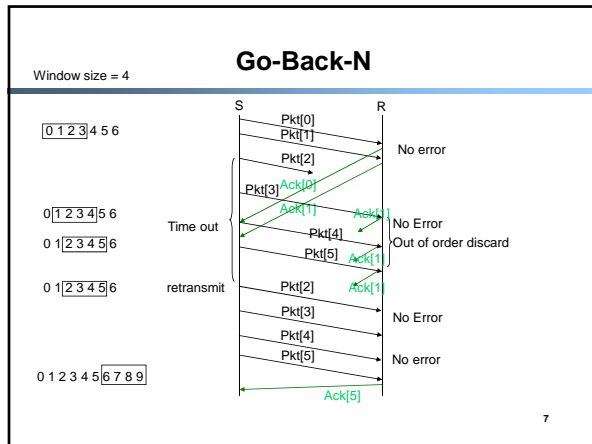
- In order



- Out of order

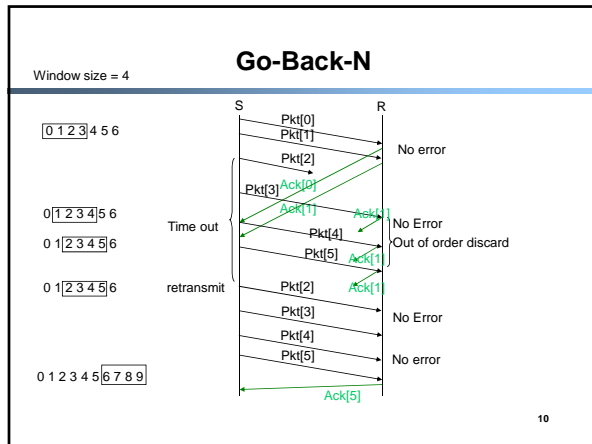


6



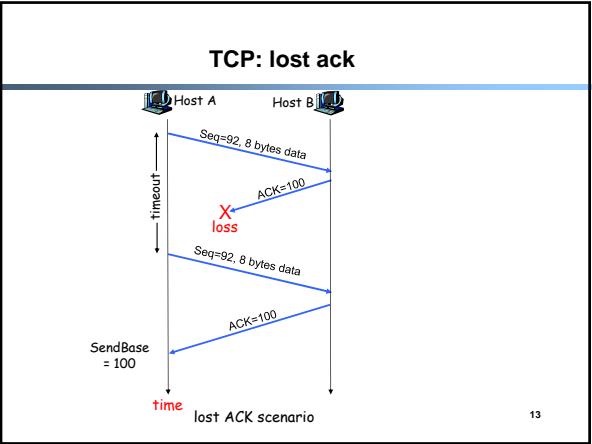
- ### Sender Actions
- Application data received
 - Wait if window is full
 - Construct segment \leftarrow nextseqnum
 - Buffer a copy
 - Set up a timer
 - Send: hand to IP layer
 - nextseqnum ++
 - ACK received
 - If $ACK(n) > send_base$: update send_base
 - Slide window
 - Otherwise, ignore
 - Timer times out
 - $timeout(n)$: retransmit packet **n** and all higher seq # packets that were sent in window
- 8

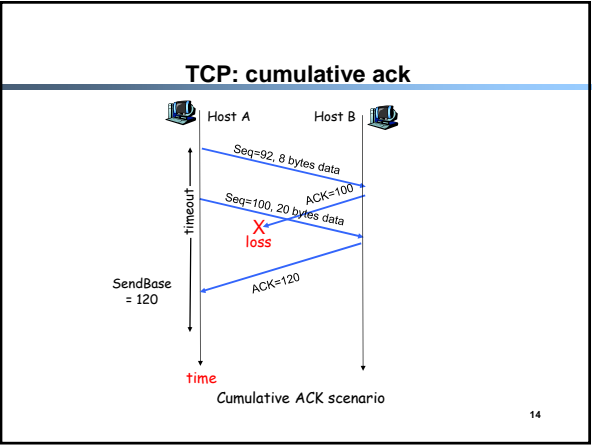
- ### Receiver Actions
- Receiver maintain: next in-order sequence number **n** (next expected segment)
 - Segment **n** received in order
 - Deliver to application layer
 - Send $ACK(n)$
 - Update next in-order sequence number
 - Segment **i** received out of order
 - Discard **i**
 - Send $ACK(\text{last correctly received seq \#})$
- 9

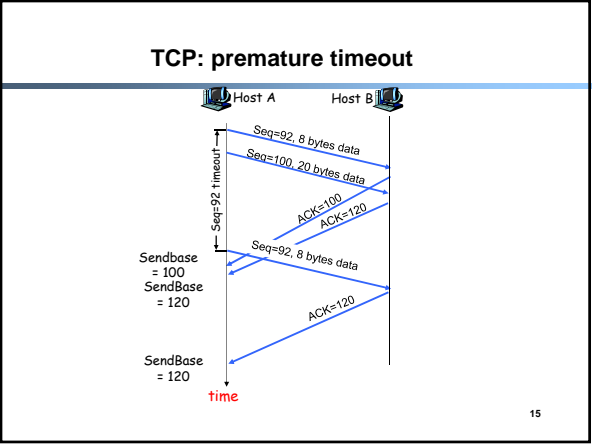


- ### Drawbacks of Go-Back-N
- Retransmit all packets that were sent but not yet acked in the window upon a time-out
 - Receiver discard out of order packets
- 11

- ### TCP reliable data transfer
- TCP creates reliable service on top of IP's unreliable service
 - Pipelined segment transmission
 - Cumulative acks
 - TCP uses single retransmission timer
 - Retransmissions are triggered by:
 - timeout events
- 12







TCP sender events:

data received from application:

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running (timer is for oldest unacked segment)

timeout:

- retransmit segment that caused timeout
- restart timer

Ack rcvd:

- If acknowledges previously unacked segments
 - update what is known to be acked
 - start timer if there are outstanding segments

16

NextSeqNum = InitialSeqNum
SendBase = InitialSeqNum

loop (forever) {
 switch(event) {

 event: data received from application above
 create TCP segment with sequence number NextSeqNum
 if (timer currently not running)
 start timer
 pass segment to IP
 NextSeqNum = NextSeqNum + length(data)

 event: timer timeout
 retransmit not-yet-acknowledged segment with
 smallest sequence number
 start timer with timeout interval doubled

 event: ACK received, with ACK field value of y
 if (y > SendBase) {
 SendBase = y
 if (there are currently not-yet-acknowledged segments)
 start timer
 }
 }
} /* end of loop forever */

TCP sender (simplified)

Comment:

- SendBase-1: last cumulatively ack'd byte

Example:

- SendBase-1 = 71; y= 73, so the rcvr wants 73+ ; y > SendBase, so that new data is acked

17


TCP ACK generation [RFC 1122, RFC 2581]

Event at Receiver	TCP Receiver action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expect seq. # . Gap detected	Immediately send duplicate ACK, indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap

18

Flow control: Window Size and Throughput

- Sliding-window based flow control:
 - larger window \rightarrow higher throughput
 - throughput = wnd/RTT
 - need to worry about sequence number wrapping
 - cumulative ack
 - timeout, retransmission
- Remember: window size controls throughput



19

Why do You Care About Congestion Control?

- Otherwise you get to congestion collapse
- How might this happen?
 - assume network is congested (a router drops packets)
 - you learn the receiver didn't get the packet
 - either by ACK, NACK, or Timeout
 - what do you do? retransmit packet
 - still receiver didn't get the packet
 - retransmit again
 - and so on ...
 - and now assume that everyone is doing the same!
- Network will become more and more congested
 - and this with duplicate packets rather than new packets!

20

Solutions?

- Increase buffer size. Why not?
- Slow down
 - if you know that your packets are not delivered because network congestion, slow down
- Questions:
 - how do you detect network congestion?
 - by how much do you slow down?


21

What's Really Happening?

- knee – point after which
 - throughput **increases very slow**
 - delay **increases fast**
- cliff – point after which
 - throughput starts to **decrease very fast to zero** (congestion collapse)
 - delay **approaches infinity**

- Note (in an M/M/1 queue)


- delay = $\frac{1/\mu}{1 - \frac{\lambda}{\mu}}$



22

Congestion Control vs. Congestion Avoidance

- Congestion control goal
 - stay left of cliff
- Congestion avoidance goal
 - stay left of knee



23

Goals

- Operate near the knee point
- Remain in equilibrium
- How to maintain equilibrium?
 - don't put a packet into network until another packet leaves. How do you do it?
 - use ACK: send a new packet only after you receive and ACK. Why?
 - maintain number of packets in network "constant"

24

How Do You Do It?

- Detect when network approaches/reaches knee point
- Stay there
- Questions
 - how do you get there?
 - what if you overshoot (i.e., go over knee point) ?
- Possible solution:
 - increase window size until you notice congestion
 - decrease window size if network congested

25

Possible Choices

- Window increase, decrease algorithms
 - additive increase, multiplicative decrease
 - multiplicative increase, additive decrease
 - multiplicative increase, multiplicative decrease
- Convergence
 - Which converge?
 - Which converge to an efficient operating point?

26

The Choice

- Additive increase, multiplicative decrease
- Intuition:
 - Let L_i denote average queue length and W_i denote window size over period I
 - on equilibrium: $L_{i+1} = N$
 - on congestion: $L_{i+1} = N + \gamma L_i \rightarrow L_{i+k} = N \sum_{l=0}^{k-1} \gamma^l + \gamma^k L_i$
 - queue size increase exponentially \rightarrow need to reduce window size at least as fast: $W_{i+1} = d W_i \quad (d < 1)$
 - on no congestion \rightarrow increase window size: $W_{i+1} = W_i + u \quad (u \ll W_{\max})$ (see [CJ89])

27

TCP Congestion Control

- Maintains three variables:
 - *cwnd* – congestion window
 - *flow_win* – flow window; receiver advertised window
 - *ssthresh* – threshold size (used to update *cwnd*)
- For sending use: $\text{win} = \min(\text{flow_win}, \text{cwnd})$

28

TCP: Slow Start

- Goal: discover congestion quickly
- How?
 - quickly increase *cwnd* until network congested → get a rough estimate of the optimal of *cwnd*
- How do we know when network is congested?
 - packet loss (TCP, [Jac88])
 - over the cliff here → congestion control
 - congestion notification (DEC Bit scheme, [RJ88])
 - over the knee but before the cliff → congestion avoidance
- How do we know a packet is lost? (latter...)

29


TCP: Slow Start

- Whenever starting traffic on a new connection, or whenever increasing traffic after congestion was experienced:
 - Set *cwnd* = 1
 - Each time a segment is acknowledged increment *cwnd* by one (*cwnd*++).
- Does Slow Start increment slowly? Not really. In fact, the increase of *cwnd* is exponential

30

Slow Start Example

- The congestion window size grows very rapidly
- TCP slows down the increase of *cwnd* when ***cwnd* ≥ *ssthresh***



31

Congestion Avoidance


- Goal: maintain operating point at the left of the cliff:
- How?
 - additive increase**: starting from the rough estimate, slowly increase *cwnd* to probe for additional available bandwidth
 - multiplicative decrease**: cut congestion window size aggressively if a timeout occurs

32

Congestion Avoidance

- Slow down "Slow Start"
- If *cwnd* > *ssthresh* then
 - each time a segment is acknowledged increment *cwnd* by $1/cwnd$ ($cwnd += 1/cwnd$).
- So *cwnd* is increased by one only if all segments have been acknowledged.
- (more about *ssthresh* latter)

33



Putting Everything Together: TCP Pseudocode

Initially:
 $CongWin = 1;$
 $ssthresh = \text{infinite};$


New ack received:
 if ($CongWin < ssthresh$)
 /* Slow Start */
 $CongWin = CongWin + 1;$
 else
 /* Congestion Avoidance */
 $CongWin = CongWin + 1 / CongWin;$

Timeout:
 /* Multiplicative decrease */
 $ssthresh = win / 2;$
 $CongWin = 1;$

3 Duplicate Acks:
 $ssthresh = win / 2;$
 $CongWin = ssthresh;$

while ($next < unack + win$)
 transmit next packet;

where $win = \min(CongWin, flow_win);$




Packet Loss Detection

- Wait for Retransmission Time Out (RTO)
- What's the problem with this?
- **Because RTO is performance killer**
- In BSD TCP implementation, RTO is usually more than 1 second
 - the granularity of RTT estimate is 500 ms
 - retransmission timeout is at least two times of RTT
- Solution: Don't wait for RTO to expire

37

Fast Retransmit

- Resend a segment after 3 duplicate ACKs
 - remember a duplicate ACK means that an out-of-sequence segment was received



- Notes:
 - duplicate ACKs due to packet reordering!
 - if window is small don't get duplicate ACKs!


38

Fast Recovery

- After a fast-retransmit set *cwnd* to *ssthresh/2*
 - i.e., don't reset *cwnd* to 1
- But when RTO expires still do *cwnd* = 1
- Fast Retransmit and Fast Recovery → implemented by TCP Reno; most widely used version of TCP today

39

Fast Retransmit and Fast Recovery



- Retransmit after 3 duplicated acks
 - prevent expensive timeouts
- No need to slow start again
- At steady state, *cwnd* oscillates around the optimal window size.

40

- TCP
 - Round trip time estimation

41

TCP Round Trip Time and Timeout

Q: how to set TCP timeout value?

- longer than RTT
- too short: premature timeout
 - unnecessary retransmissions
- too long: slow reaction to segment loss

Q: how to estimate RTT?

- **SampleRTT**: measured time from segment transmission until ACK receipt
 - ignore retransmissions
- **SampleRTT** will vary, want estimated RTT "smoother"
 - average several recent measurements, not just current **SampleRTT**

42


TCP Round Trip Time and Timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- Exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.125$

43

Example RTT estimation:



44

TCP Round Trip Time and Timeout

Setting the timeout

- **EstimatedRTT** plus "safety margin"
 - large variation in **EstimatedRTT** → larger safety margin
- first estimate of how much **SampleRTT** deviates from **EstimatedRTT**:

$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)

Then set timeout interval:

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

45

Congestion Control Summary

- Architecture: end system detects congestion and slow down
- Starting point:
 - slow start/congestion avoidance
 - packet drop detected by retransmission timeout RTO as congestion signal
 - fast retransmission/fast recovery
 - packet drop detected by three duplicate acks
- Router support
 - Binary feedback scheme: explicit signaling
 - Today Explicit Congestion Notification [RFC99]

46

Reflection

- TCP implicitly assumes that **all** sources need to cooperate
- What are implications?

47
